

## Module 04 – Project


# SALES PREDICTION

Nguyen Quoc Thai

# Objectives

## Regression

- ❖ Regression Task
- ❖ Linear Regression
- ❖ Non-linear Regression
- ❖ Using Sklearn library



	TV	Radio	Social Media	Influencer	Sales
0	16.0	6.566231	2.907983	Mega	54.732757
1	13.0	9.237765	2.409567	Mega	46.677897
2	41.0	15.886446	2.913410	Mega	150.177829
3	83.0	30.020028	6.922304	Mega	298.246340
4	15.0	8.437408	1.405998	Micro	56.594181

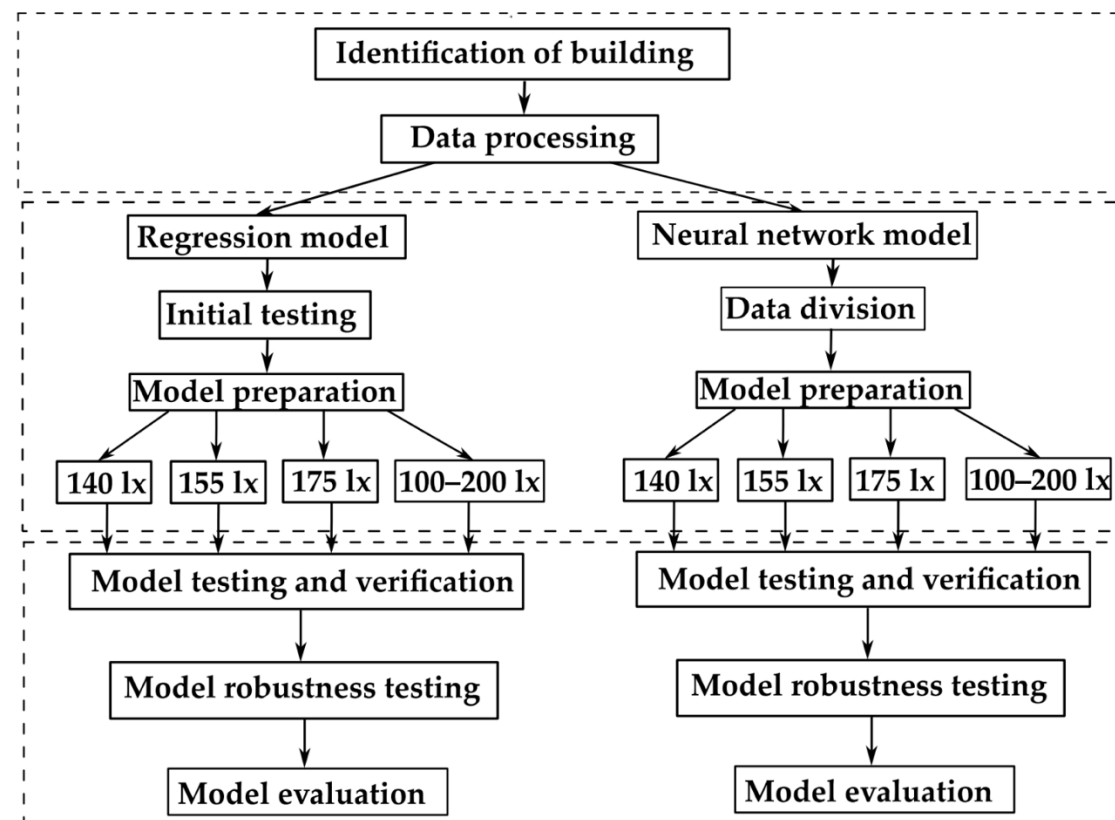
## Sales Prediction

- ❖ Exploratory Data Analysis (EDA)
- ❖ Feature Scaling
- ❖ Modeling
- ❖ Evaluation
- ❖ Custom Polynomial Features



## Research ideas for polynomial regression applications


[A Comparative Analysis of Polynomial Regression and Artificial Neural Networks for Prediction of Lighting Consumption](#)



Các nhóm (Gồm các thành viên thuộc AIO) có lời giải sơ bộ, AD Vinh sẽ hướng dẫn tiếp để ra paper.

## SECTION 1

## Regression



## SECTION 2

## Sales Prediction

	TV	Radio	Social Media	Influencer	Sales
0	16.0	6.566231	2.907983	Mega	54.732757
1	13.0	9.237765	2.409567	Mega	46.677897
2	41.0	15.886446	2.913410	Mega	150.177829
3	83.0	30.020028	6.922304	Mega	298.246340
4	15.0	8.437408	1.405998	Micro	56.594181

# Regression



## Regression Task

### Regression



- Predict a continuous value based on the input variables




What will be the temperature tomorrow?



### Training Data



### Test Data



# Regression




## Regression Task

**Data**

Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47

Level	Salary
3,5	?
10	?

Learning



Prediction

# Regression



## Linear Regression

### Data


Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47

### Modeling

$$y = wx + b$$

Find  $w$  and  $b$  to fit the data

### Visualization



# Regression



## Linear Regression using Gradient Descent

**Data**

Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47

**Inputs / Features**

$$X = \begin{bmatrix} 1 & \varphi_1(1) & \dots & \varphi_{d-1}(1) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \varphi_1(N) & \dots & \varphi_{d-1}(N) \end{bmatrix}$$

**Target**

$$Y = \begin{bmatrix} y(1) \\ \vdots \\ y(N) \end{bmatrix}$$

**Weight**

$$\theta = \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_{d-1} \end{bmatrix}$$

**Predict**

$$\hat{Y} = \begin{bmatrix} \theta_0 + \theta_1 * \varphi_1(1) + \dots + \theta_{d-1} \varphi_{d-1}(1) \\ \vdots \\ \theta_0 + \theta_1 * \varphi_1(N) + \dots + \theta_{d-1} \varphi_{d-1}(N) \end{bmatrix}$$




# Regression



## Linear Regression using Gradient Descent

### Data



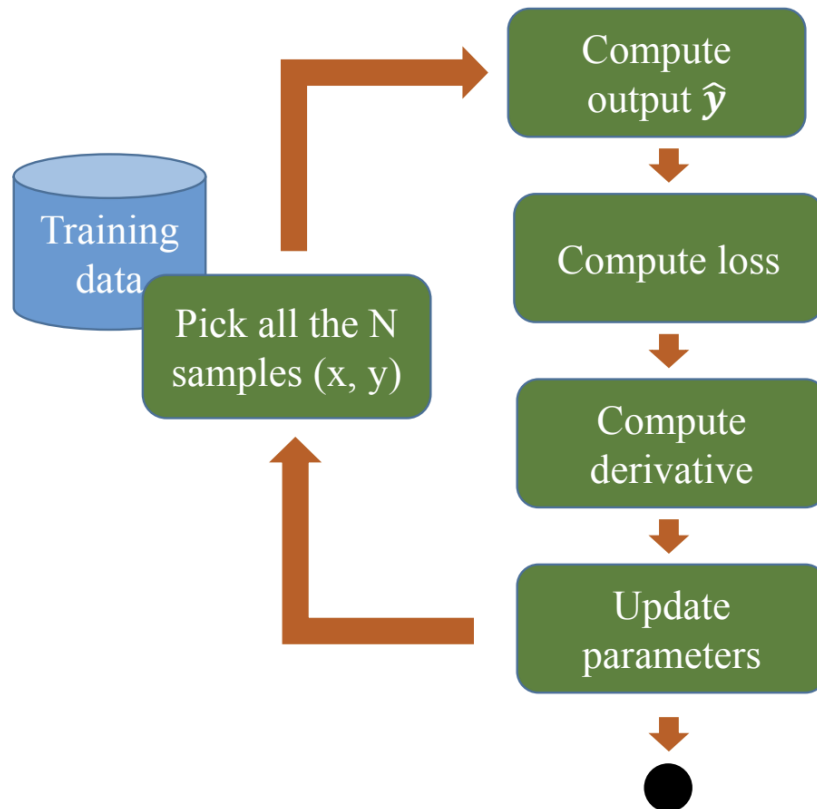
$$X = \begin{bmatrix} 1 & \varphi_1(1) & \dots & \varphi_{d-1}(1) \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \varphi_1(N) & \dots & \varphi_{d-1}(N) \end{bmatrix}$$

$$Y = \begin{bmatrix} y(1) \\ \vdots \\ y(N) \end{bmatrix}$$

$$y = wX + b$$

$$\theta = \begin{bmatrix} b \\ w \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_{d-1} \end{bmatrix}$$

### Learning



1) Pick all the N samples from training data

2) Compute output  $\hat{y}$

$$\hat{y} = x\theta$$

3) Compute loss

$$L(\hat{y}, y) = (\hat{y} - y) \odot (\hat{y} - y)$$

4) Compute derivative

$$k = 2(\hat{y} - y)$$

$$L'_{\theta} = x^T k$$

5) Update parameters

$$\theta = \theta - \eta \frac{L'_{\theta}}{N} \quad \eta \text{ is learning rate}$$

# Regression



## Limitation

Modeling

$$y = wx + b$$


The **main disadvantage** of this technique is that **the model is linear in both the parameters and the features**. This is a very restrictive assumption, quite of the often **data exhibits behaviours that are nonlinear in the features**


Extend this approach to more flexible models...

# Regression



## Moving Beyond Linearity

Data Visualization



Linear function

$$\hat{y}(i) = \theta_0 + \theta_1 * \varphi(i)$$

Polynomial function

$$\hat{y}(i) = \theta_0 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2$$

Nonlinear regression estimates the output based on nonlinear function


Notice that the prediction is **still linear in the parameters** but **nonlinear in the features**

# Regression




## Moving Beyond Linearity


2-degree polynomial function




3-degree polynomial function




Step function



Sinusoidal



Exponential function



# Regression




## Polynomial Regression

2-degree polynomial function

$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2$$

Find  $\theta_0$  ,  $\theta_1$  ,  $\theta_2$  to fit the data



# Regression



## Polynomial Features

2-degree polynomial function

$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2$$

Create polynomial feature



$\psi(\cdot)$  is referred to as basis function and it can be seen as a function that transforms the input in some way (In this case its powers function)

# Regression

## ! Polynomial Features

2-degree polynomial function

$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2$$

Create polynomial feature

**Data**

Level	Salary
0	45000
1	50000
2	60000
3	80000
4	110000
5	160000

**Input**

0
1
2
3
4
5

**1**

1
1
1
1
1
1

**$\psi(\varphi(i))$**

**$\varphi(i)$**

0
1
2
3
4
5

**$\varphi(i)^2$**

0
1
4
9
16
25

# Regression



## Polynomial Features

2-degree polynomial function

$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2$$

Features			Target
1	$\varphi(i)$	$\varphi(i)^2$	
1	0	0	45000
1	1	1	50000
1	2	4	60000
1	3	9	80000
1	4	16	110000
1	5	25	160000



# Regression



## Polynomial Features

3-degree polynomial function

$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2 + \theta_3 * \varphi(i)^3$$

Features				Target
1	$\varphi(i)$	$\varphi(i)^2$	$\varphi(i)^3$	
1	0	0	0	45000
1	1	1	1	50000
1	2	4	8	60000
1	3	9	27	80000
1	4	16	64	110000
1	5	25	125	160000

# Regression



## Polynomial Features

### Input

0
1
2
3
4
5

### Features

1	$\varphi(i)$	$\varphi(i)^2$
1	0	0
1	1	1
1	2	4
1	3	9
1	4	16
1	5	25

### Algorithm

```
def create_polynomial_features(X, degree=2):  
    """Creates the polynomial features  
    Args:  
        X: A torch tensor for the data.  
        degree: A integer for the degree of  
              the generated polynomial function.  
    """  
    X_new = X  
    for d in range(2, degree+1):  
        X_new = np.c_[X_new, np.power(X, d)]  
    return X_new
```

# Regression



## Polynomial Features

### Input

0
1
2
3
4
5

### Features

1	$\varphi(i)$	$\varphi(i)^2$
1	0	0
1	1	1
1	2	4
1	3	9
1	4	16
1	5	25

```
1 from sklearn.preprocessing import PolynomialFeatures
```

```
1 poly_features = PolynomialFeatures(degree=2)
```

```
1 X.to_frame()
```

```
(10, 1)
```

```
1 X_poly = poly_features.fit_transform(X.to_frame())
2 X_poly
```

```
array([[ 1.,  1.,  1.],
       [ 1.,  2.,  4.],
       [ 1.,  3.,  9.],
       [ 1.,  4., 16.],
       [ 1.,  5., 25.],
       [ 1.,  6., 36.],
       [ 1.,  7., 49.],
       [ 1.,  8., 64.],
       [ 1.,  9., 81.],
       [ 1., 10., 100.]])
```

# Regression



## Vectorization

**Data**

Level	Salary
0	45000
1	50000
2	60000
3	80000
4	110000
5	160000

**Inputs / Features with b-degree**

$$X = \begin{bmatrix} 1 & \varphi_1(1) & \dots & \varphi_1(1)^b \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \varphi_1(N) & \dots & \varphi_1(N)^b \end{bmatrix}$$

**Target**

$$Y = \begin{bmatrix} y(1) \\ \vdots \\ y(N) \end{bmatrix}$$

**Weight**

$$\theta = \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_b \end{bmatrix}$$

**Predict**

$$\hat{Y} = \begin{bmatrix} \theta_0 + \theta_1 * \varphi_1(1) + \dots + \theta_b \varphi_1(1)^b \\ \vdots \\ \theta_0 + \theta_1 * \varphi_1(N) + \dots + \theta_b \varphi_1(N)^b \end{bmatrix}$$

# Regression



## Model

### Nonlinear Regression Model

$$X = \begin{bmatrix} 1 & \varphi_1(1) & \dots & \varphi_1(1)^b \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \varphi_1(N) & \dots & \varphi_1(N)^b \end{bmatrix}$$

$$Y = \begin{bmatrix} y(1) \\ \vdots \\ y(N) \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_b \end{bmatrix}$$

$$\hat{Y} = \begin{bmatrix} \theta_0 + \theta_1 * \varphi_1(1) + \dots + \theta_b \varphi_1(1)^b \\ \vdots \\ \theta_0 + \theta_1 * \varphi_1(N) + \dots + \theta_b \varphi_1(N)^b \end{bmatrix}$$

Both models are linear  
in the parameters

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N (\hat{y} - y)^2$$

Using Gradient Decent

### Linear Regression Model

$$X = \begin{bmatrix} 1 & \varphi_1(1) \\ \vdots & \vdots \\ 1 & \varphi_1(N) \end{bmatrix}$$

$$Y = \begin{bmatrix} y(1) \\ \vdots \\ y(N) \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix}$$

$$\hat{Y} = \begin{bmatrix} \theta_0 + \theta_1 * \varphi_1(1) \\ \vdots \\ \theta_0 + \theta_1 * \varphi_1(N) \end{bmatrix}$$

# Regression

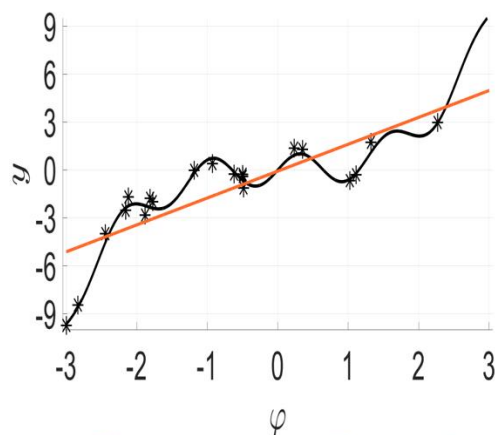


## Degree Choice

b-degree polynomial function

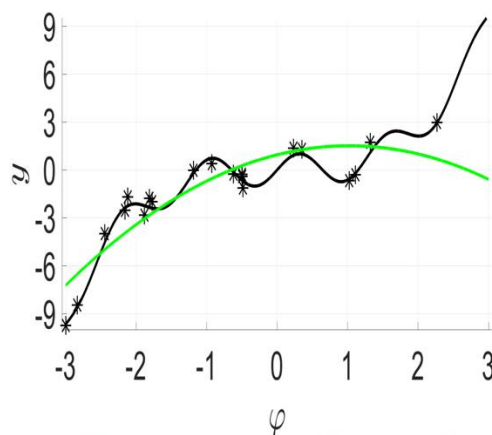
$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2 + \dots + \theta_b * \varphi(i)^b$$

The choice of the degree of the polynomial is critical and depends on the dataset at hand




**1-degree**

Too simple/not flexible enough




**2-degree**



**3-degree**

Just right



**9-degree**

Overfitting

# Regression



## Degree Choice

b-degree polynomial function

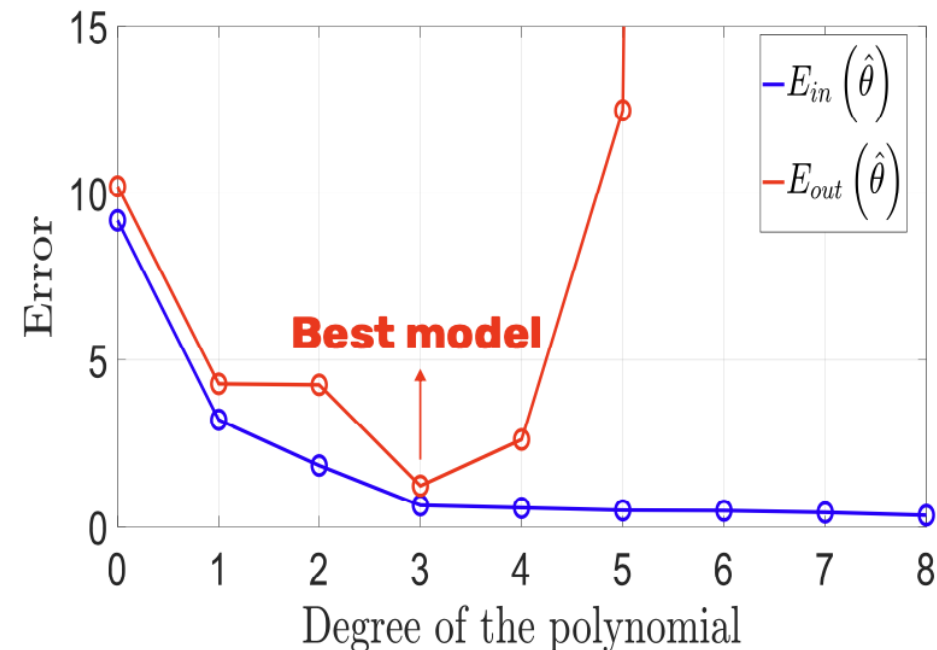
$$\hat{y}(i) = \theta_0 * 1 + \theta_1 * \varphi(i) + \theta_2 * \varphi(i)^2 + \dots + \theta_b * \varphi(i)^b$$

The choice of the degree of the polynomial is critical and depends on the dataset at hand

Good method for choice of the degree:

K-fold cross-validation




Choose the degree which has the lowest out-of-sample error



# Regression

## ! Disadvantages

**Increasing the degree of the polynomial always results in a model that is more sensitive to stochastic noise** (even if that degree is the best one obtained from validation), especially at the boundaries (where we often have less data).






QUIZ TIME

## SECTION 1

## Regression



## SECTION 2

## Sales Prediction

	TV	Radio	Social Media	Influencer	Sales
0	16.0	6.566231	2.907983	Mega	54.732757
1	13.0	9.237765	2.409567	Mega	46.677897
2	41.0	15.886446	2.913410	Mega	150.177829
3	83.0	30.020028	6.922304	Mega	298.246340
4	15.0	8.437408	1.405998	Micro	56.594181

# Sales Prediction



## Dataset




**Dataset**



**Preprocessing**

**Exploratory  
Data Analysis**



**Feature Scaling**



**Modeling**



**Evaluation**

# Sales Prediction



## Dataset

```
1 import pandas as pd
2
3 df = pd.read_csv('./SalesPrediction.csv')
4 df
```

	TV	Radio	Social Media	Influencer	Sales
0	16.0	6.566231	2.907983	Mega	54.732757
1	13.0	9.237765	2.409567	Mega	46.677897
2	41.0	15.886446	2.913410	Mega	150.177829
3	83.0	30.020028	6.922304	Mega	298.246340
4	15.0	8.437408	1.405998	Micro	56.594181



# Sales Prediction



## Dataset

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4572 entries, 0 to 4571  
Data columns (total 5 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0    TV              4562 non-null   float64  
1    Radio            4568 non-null   float64  
2    Social Media     4566 non-null   float64  
3    Influencer       4572 non-null   object  
4    Sales            4566 non-null   float64  
dtypes: float64(4), object(1)  
memory usage: 178.7+ KB
```

1 df.describe()

	TV	Radio	Social Media	Sales
count	4562.000000	4568.000000	4566.000000	4566.000000
mean	54.066857	18.160356	3.323956	192.466602
std	26.125054	9.676958	2.212670	93.133092
min	10.000000	0.000684	0.000031	31.199409
25%	32.000000	10.525957	1.527849	112.322882
50%	53.000000	17.859513	3.055565	189.231172
75%	77.000000	25.649730	4.807558	272.507922
max	100.000000	48.871161	13.981662	364.079751





## Preprocessing – One hot encoding

```
1 df = pd.get_dummies(df)
2 df
```

Radio	Social Media	Sales	Influencer_Macro	Influencer_Mega	Influencer_Micro	Influencer_Nano
6.566231	2.907983	54.732757	False	True	False	False
9.237765	2.409567	46.677897	False	True	False	False
15.886446	2.913410	150.177829	False	True	False	False
30.020028	6.922304	298.246340	False	True	False	False
8.437408	1.405998	56.594181	False	False	True	False



## Preprocessing – Handling Missing Values

```
1 df.isnull().sum()
```

	0
TV	10
Radio	4
Social Media	6
Sales	6
Influencer_Macro	0
Influencer_Mega	0
Influencer_Micro	0
Influencer_Nano	0

dtype: int64

```
1 df = df.fillna(0)
2 df.isnull().sum()
```

	0
TV	0
Radio	0
Social Media	0
Sales	0
Influencer_Macro	0
Influencer_Mega	0
Influencer_Micro	0
Influencer_Nano	0

dtype: int64

```
1 df = df.fillna(df.mean())
2 df.isnull().sum()
```

	0
TV	0
Radio	0
Social Media	0
Sales	0
Influencer_Macro	0
Influencer_Mega	0
Influencer_Micro	0
Influencer_Nano	0



dtype: int64



# Sales Prediction



## EDA






# Sales Prediction



## EDA

```
1 df[['TV', 'Radio', 'Social Media', 'Sales']].corr()
```

	TV	Radio	Social Media	Sales
TV	1.000000	0.860518	0.522565	0.988570
Radio	0.860518	1.000000	0.604450	0.863790
Social Media	0.522565	0.604450	1.000000	0.526777
Sales	0.988570	0.863790	0.526777	1.000000



# Sales Prediction



## Train Test Split

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(
3     X,
4     y,
5     test_size=0.33,
6     random_state=0
7 )
```

```
1 X_train.shape, X_test.shape
```

```
((3063, 7), (1509, 7))
```

```
1 y_train.shape, y_test.shape
```

```
((3063, 1), (1509, 1))
```



## Feature Scaling

### MaxAbsScaler

$$x_{new} = \frac{x}{x_{max}}$$

### MinMaxScaler

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

### StandardScaler

$$x_{new} = \frac{x - \mu}{\sigma}$$

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4 X_train_processed = scaler.fit_transform(X_train)
```

```
1 scaler.mean_
array([53.9970617 , 18.22209011,  3.33487105,  0.24779628,  0.25138753,
        0.25008162,  0.25073457])
```

```
1 scaler.scale_
array([26.24285095,  9.6336957 ,  2.21929717,  0.43173288,  0.43381083,
        0.43305981,  0.43343598])
```

```
1 X_test_processed = scaler.transform(X_test)
```

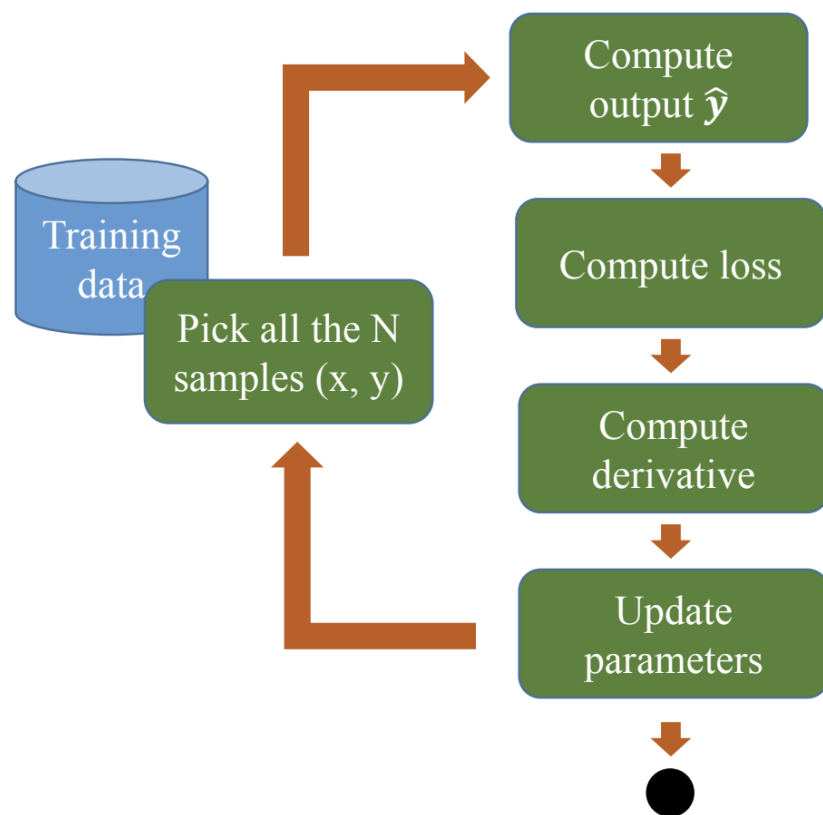
```
1 X_train_processed.shape
(3063, 7)
```

```
1 X_test_processed.shape
(1509, 7)
```

# Sales Prediction



## Modeling



1) Pick all the N samples from training data

2) Compute output  $\hat{y}$

$$\hat{y} = x\theta$$

3) Compute loss

$$L(\hat{y}, y) = (\hat{y} - y) \odot (\hat{y} - y)$$

4) Compute derivative

$$k = 2(\hat{y} - y)$$

$$L'_{\theta} = x^T k$$

5) Update parameters

$$\theta = \theta - \eta \frac{L'_{\theta}}{N} \quad \eta \text{ is learning rate}$$



# Sales Prediction



## Modeling

```
1 from sklearn.linear_model import LinearRegression
2
3 linear_model = LinearRegression()
4 linear_model.fit(X_train_processed, y_train)
```

▼ LinearRegression ⓘ ?

```
LinearRegression()
```

```
1 preds = linear_model.predict(X_test_processed)
2 r2_score(y_test, preds)
```

0.9820910569999272

# Sales Prediction



## Polynomial Regression

```
1 from sklearn.preprocessing import PolynomialFeatures
2
3 poly_features = PolynomialFeatures(degree=2)
4 X_train_poly = poly_features.fit_transform(X_train_processed)
```

```
1 X_train_poly
```

```
array([[ 1.          ,  0.34306251, -0.39269809, ...,  2.99869452,
        -1.00174084,  0.33464052],
       [ 1.          , -0.19041611, -0.28821416, ...,  0.33347845,
        0.33405898,  0.33464052],
       [ 1.          , -0.41904981, -1.07312224, ...,  0.33347845,
        -0.99826219,  2.98828125],
       ...,
       [ 1.          , -1.6003239 , -1.72760008, ...,  0.33347845,
        -0.99826219,  2.98828125],
       [ 1.          , -0.57147227, -0.9126861 , ...,  2.99869452,
        -1.00174084,  0.33464052],
       [ 1.          , -1.25737336, -1.45632493, ...,  2.99869452,
        -1.00174084,  0.33464052]])
```

```
1 X_test_poly = poly_features.transform(X_test_processed)
```

# Sales Prediction



## Polynomial Regression

```
1 poly_model = LinearRegression()  
2 poly_model.fit(X_train_poly, y_train)
```

▼ LinearRegression ⓘ ?

LinearRegression()


```
1 preds = poly_model.predict(X_test_poly)  
2 r2_score(y_test, preds)
```

0.9785743009106321

# Sales Prediction



## Custom Polynomial Regression



```
1 X_train_processed[:, 2:3]
```

```
array([[ -0.17117575],
       [ -1.47454833],
       [ -0.55726535],
       ...,
       [  0.58703816],
       [ -1.22457248],
       [ -1.04684805]])
```

```
1 x_train_poly = create_polynomial_features(X_train_processed[:, 2:3], degree=2)
```

```
2 x_train_poly
```

```
array([[ -0.17117575,  0.02930114],
       [ -1.47454833,  2.17429276],
       [ -0.55726535,  0.31054467],
       ...,
       [  0.58703816,  0.3446138 ],
       [ -1.22457248,  1.49957777],
       [ -1.04684805,  1.09589085]])
```



# Sales Prediction



## Custom Polynomial Regression

```
1 x_test_poly = create_polynomial_features(X_test_processed[:, 2:3], degree=2)
2 X_test_poly = np.hstack((X_test_processed, x_test_poly[:, 1:]))
3 X_test_poly
```

```
array([[ -0.34283858, -0.11361891, -0.84351677, ...,  1.73167391,
        -0.57848122,  0.71152054],
       [  0.76222428,  1.17276695, -0.45136527, ..., -0.57747593,
        1.72866459,  0.20373061],
       [  1.14328044,  1.04152694,  1.06570814, ..., -0.57747593,
        -0.57848122,  1.13573384],
       ...,
       [  0.95275236,  1.11773825,  1.06866838, ..., -0.57747593,
        -0.57848122,  1.14205211],
       [-0.41904981, -0.32445517,  1.20063234, ..., -0.57747593,
        -0.57848122,  1.44151802],
       [-0.91442282, -1.25598448, -0.32806086, ..., -0.57747593,
        -0.57848122,  0.10762393]])
```

```
1 X_train_poly.shape, X_test_poly.shape
```

```
((3063, 8), (1509, 8))
```

```
1 poly_model = LinearRegression()
2 poly_model.fit(X_train_poly, y_train)
```

▼ LinearRegression ⓘ ?

```
LinearRegression()
```

```
1 preds = poly_model.predict(X_test_poly)
2 r2_score(y_test, preds)
```

```
0.9820873203866817
```



# Sales Prediction




## Results

Model	R2
Linear Regression	0.9821
Polynomial Regression (2)	0.9786
Custom Polynomial Regression	0.9821

# Summary

## Regression

- ❖ Regression Task
- ❖ Linear Regression
- ❖ Non-linear Regression
- ❖ Using Sklearn library



	TV	Radio	Social Media	Influencer	Sales
0	16.0	6.566231	2.907983	Mega	54.732757
1	13.0	9.237765	2.409567	Mega	46.677897
2	41.0	15.886446	2.913410	Mega	150.177829
3	83.0	30.020028	6.922304	Mega	298.246340
4	15.0	8.437408	1.405998	Micro	56.594181

## Sales Prediction

- ❖ Exploratory Data Analysis (EDA)
- ❖ Feature Scaling
- ❖ Modeling
- ❖ Evaluation
- ❖ Custom Polynomial Features



AI VIET NAM

@aivietnam.edu.vn

# Thanks!

## Any questions?